
DEEP IMPLICIT LAYERS FOR SMOOTH IMPLIED VOLATILITY SURFACE INTERPOLATION

 **Rafael Zimmer**

Institute of Mathematical Science and Computing
University of São Paulo - Brazil
rafael.zimmer@usp.br

 **Adalton Sena Filho**

Institute of Mathematical Science and Computing
University of São Paulo - Brazil
adaltonsena@usp.br

 **Luís Eduardo de Brito Câmara**

Institute of Mathematical Science and Computing
University of São Paulo - Brazil
eduardobritoc@usp.br

15th of July, 2024

ABSTRACT

In this paper, we implement a neural network architecture based on deep implicit layers with the objective of optimizing the parameters of multiple SABR (stochastic alpha, beta, rho) models. The parameter values for the models are updated daily and used for generating smooth implied volatility surfaces in the context of financial options. The motivation behind our approach is to enable smooth interpolation and extrapolation of the surface, i.e., making it infinitely differentiable at all points while still maintaining some level of interpretability within our model when compared to state-of-the-art deep transformer architectures. The main contribution of the proposed architecture is to solve the problem of selecting the candidates α, ρ, ν to reduce statistical arbitrage opportunities for the respective generated surface. Our methodology includes using historical observed values of listed options as well as the calculated implied volatility for these contracts. The innate ability of transformer architectures to process observations of sequences with non-fixed lengths, i.e., for contracts with varying expiration dates is used to choose candidate points that best replicate the daily surfaces. Using our approach, the resulting surface curvature and volatility values are then compared to baseline classical methods such as linear interpolations, splines and non-implicit transformer networks.

Keywords Neural Networks · Implied Volatility · Options · Quantitative Methods · Implicit Layers

1 Introduction

1.1 A Brief Introduction to Implied Volatility and Surface Reconstruction

Option contracts are one of the most important types of derivative contracts, and implied volatility surfaces play a critical role in the context of finding the correct no-arbitrage price [Ludkovski, 2023] for these contracts. A derivative is an asset traded with respect to another asset (also called the underlying asset). An option specifically allows the buyer of the option to buy or sell (also called call and put contracts) a specified asset from the seller at a predefined future date. The pricing equations of these derivatives called the Black-Scholes model depend on the time to maturity of the contract, the price of the underlying asset to be traded, the price of the asset to be traded and the risk-free rate [Black and Scholes, 1973]. There is also an additional variable for contract pricing, called the implied volatility (IV or σ_{iv}). This value refers to the volatility considered, or to be expected, for the asset price at the time of maturity, and is an essential variable in the Black-Scholes equation for pricing option contracts.

Having access to the true implied volatility values allows investors to make sales at the correct prices consistent with the current no-arbitrage price of the contracts within the market. However, it is not always possible to obtain the true IV values that fully reflect the characteristics of existing offers. Ensuring a way to obtain this value with confidence and readily available data in the market is a crucial task for quantitative analysts [Kim et al., 2006]. Traditional strategies for daily estimation of these values are predominantly based on linear or spline interpolation algorithms [Homescu, 2011] or stochastic differential equation models [Jordan and Tier, 2011] that use numerical methods to approximate their values from existing option contracts in the market. Both methods try to approximate the correct value for the implied volatility for each pair of time to maturity and future asset price currently listed (also called strike price). This approach allows visualizing the model results in the form of a surface, where the $\mathbf{XY}$ axes are the maturity and strikes pairs and the $\mathbf{Z}$ axis is the $IVOL$. This surface is called the **implied volatility surface** (IVS). The surface, noted as $\mathcal{S}$ is therefore a function of the maturity time t and the strike price K ,

$$\mathcal{S} = f(t, K), \text{ where } f : \mathbb{R}^2 \rightarrow \mathbb{R} \quad (1)$$

Ensuring accurate predictions of the IVS is crucial, as even small errors can lead to incorrect pricing and consequently large losses of profit or even negative return [Hagan et al., 2014] given the leveraging capabilities of options when compared to non-derivative securities. Classic models such as linear interpolations, polynomial fitting, and *splines* generate arbitrage opportunities in the resulting IVS and often struggle to effectively capture market dynamics, besides lacking flexibility for extrapolation [Hagan et al., 2014].

1.2 Contributions of the Presented Work

In this paper, we propose an approach to address the challenges of interpolation and extrapolation of the IVS, reducing arbitrage opportunities. The chosen method uses a stochastic differential equation solution approach, specifically the *Stochastic Alpha, Beta, Rho* (*SABR*) model. Historical values of observed implied volatilities are then used to optimize the weights of a nonlinear deep neural network.

The parameters α, ρ, ν of the *SABR* model are obtained through a simple implementation of the Broyden-Fletcher-Goldfarb-Shanno (BFGS) quasi-Newton non-linear optimization algorithm [Broyden, 1970, Fletcher, 1970, Goldfarb, 1970, Shanno, 1970]. This step is commonly performed using only a subset of the points observed daily [Kim et al., 2021], due to the fact that using all points introduces various outliers that do not match the real prices of the derivative, resulting in arbitrage opportunities in the final model. To solve this difficulty, parameters are often manually filtered according to criteria set by human experts [Hagan et al., 2002], and a subset of candidate points is used for parameter optimization. Rather than manually selecting these candidate points, the use of models capable of receiving point clouds and returning a score to select the best points automatically has been proposed by [Hagan et al., 2002] instead. The main contribution of this research will be the use of *neural implicit layers*, a variation of fully connected linear layers that solve a fixed point equation within the final transformer architecture. This type of neural layers has been used efficiently for representing smooth densities and state-of-the-art performance when compared to plain *transformer* architectures [Kolter et al., 2020, Kawaguchi, 2021]. Unlike regular linear neural layers, the chosen architecture can reduce the size and complexity of the final model *Kolter2020* while still being more suitable for handling unordered point clouds over time [Kim et al., 2024] when compared to recurrent neural layers. The optimization and filtering steps are discussed in more detail under the development section.

1.3 The Transformer Architecture

The *SABR* model is a stochastic volatility model used to obtain smooth surfaces, i.e., continuous and infinitely differentiable at all points [Hagan et al., 2015]. In the context of this work, a score will be assigned for each candidate point in the daily listed options contracts. This score, called the *summarized suitability* score (*SS* score) corresponds to the suitability of the candidate point to summarize the characteristics of the generated surface, i.e., how feasible it is to reduce the number of candidates when optimizing the values for one of the parameters of the *SABR* model without altering the resulting surface shape or its smoothness¹. In summary, this score value is used to remove points with the lowest score, that may introduce arbitrage in the optimization process of the *SABR* models.

Usually, the choice of subsets and the scoring process is done manually by analysts, as mentioned earlier, and the next section will focus on the details covering the *SABR* model and how we use a implicit neural layer with attention, specifically an **encoder** architecture for assigning scores automatically given the daily points candidate set. This approach is an implicit fixed-point solver variation for the scoring model proposed in [Kim et al., 2021]

1.4 Deep Implicit Layers

Differently from fully connected layers in which the transformation from input to output is defined explicitly, implicit layers are instead defined by conditions that the output must satisfy. For explicit layers, a given output z is computed directly from the input x using a function f , such that $z = f(x)$. In contrast, an implicit layer is the function g so that it depends on both the input x and the expected output z , where z is a solution for the equation $g(x, z) = 0$.

Formulating the problem as a fixed point equation can be used for a multitude of problems, including finding fixed points in recurrent networks, solutions to differential equations leading to Neural ODEs, and several optimization problems [Chen et al., 2019]. The implicit formulation offers several practical advantages, primarily the separation of the solution method from the layer definition which can be useful in maintaining interpretability of the resulting model [Kawaguchi, 2021].

Moreover, implicit layers offer significant benefits specifically in deep learning and automatic differentiation frameworks. Automatic differentiation methods store the entire computation graph, which can be memory-intensive. With implicit layers, however, the implicit function theorem allows computing gradients directly at the solution point and therefore storing intermediate variables becomes redundant [Massaroli et al., 2021]. This improves memory efficiency, making implicit layers particularly advantageous when substituting existing layers within large deep learning models [Li et al., 2020].

¹The use of a per-point score is similar to what the principal components of a collection of points is within the Principal component analysis technique

2 Implementation of a Parametric SABR Model and Implicit Neural Network Architectures

2.1 Defining the SABR Model Equations

The SABR model is a stochastic differential model, defined by the following equations in which the dynamics of the forward price and the corresponding volatility levels are used to calculate the implied volatility.

- **Alpha** (α): This is the initial volatility level.
- **Beta** (β): The elasticity parameter between the volatility of the asset and its price.
- **Rho** (ρ): Statistical correlation between the asset price and its volatility.
- **Volvol** (ν): The volatility of the volatility process itself.

The implied volatility function of the SABR model is then given by:

$$\sigma_{BS}(f, K) = \alpha \frac{(fK)^{\frac{1-\beta}{2}}}{1 + \left(\frac{(1-\beta)^2 \rho^2}{24} + \frac{(1-\beta)^2 (2-3\rho^2)}{1920} \right) (fK)^{1-\beta}} \quad (2)$$

In the context of this paper, multiple SABR models are fitted daily, corresponding to each unique maturity value at the respective day. A continuous function is used to map the different maturities to the three model parameters (values for β are fixed and remain unchanged). The three variables used to encode the maturity levels are the following:

- **P**: Given a function $\alpha_t = f_\alpha(t, \mathbf{P}) \rightarrow \mathbb{R}$, the values for α for each individual *SABR* model are encoded by $\mathbf{P}$, such that $\mathbf{P} \in \mathbb{R}^5$.
- **Q**: Given a function $\rho_t = f_\rho(t, \mathbf{Q}) \rightarrow \mathbb{R}$, the values for ρ for each individual *SABR* model are encoded by $\mathbf{Q}$, such that $\mathbf{Q} \in \mathbb{R}^4$.
- **R**: Given a function $\nu_t = f_\nu(t, \mathbf{R}) \rightarrow \mathbb{R}$, the values for ν for each individual *SABR* model are encoded by $\mathbf{R}$, such that $\mathbf{R} \in \mathbb{R}^4$.

And finally, the daily parametric SABR model is given by the set of encoded parameters $\{P, Q, R\}$ and the following equations to calculate values for $f_\alpha, f_\rho, \text{e } f_\nu$ at each unique maturity value t :

$$f_\alpha(t, \mathbf{P}) = p_0 + \frac{p_3}{p_4} \left(\frac{1 - e^{-p_4 t}}{p_4 t} \right) + \frac{p_1}{p_2} e^{-p_2 t} \quad (3)$$

$$f_\rho(t, \mathbf{Q}) = q_0 + q_1 t + q_2 e^{-q_3 t} \quad (4)$$

$$f_\nu(t, \mathbf{R}) = r_0 + r_1 t^{r_2} e^{r_3 t} \quad (5)$$

The task of tuning the SABR model parameters is a supervised regression task. In machine learning, regression refers to problems where the output is a continuous variable, as opposed to classification, where the output is a discrete category. In our case, the output of the transformer model are the continuous SS scores for each parameter α, ρ , and ν . Then, the top $p = 80\%$ candidate points are passed as inputs to the BFGS algorithm, targeting the values $\mathbf{P}, \mathbf{Q}$, and $\mathbf{R}$ that minimize possible arbitrage opportunities of contracts with interpolated or extrapolated maturities and strikes.

2.2 Proposed Encoder Architecture and Error Functions

Transformers are a specific type of neural network architecture capable of capturing complex dependencies in sequential non-fixed length data. Within our proposed architecture, we use a transformer encoder to approximate the SS scores for each candidate point $\mathcal{S}_i^t \in \mathcal{S}^t$ from the daily observed input. Then, for each training epoch it estimates a loss score for the mean squared error between the encoder's output and a z-score obtained by Monte Carlo estimation from a candidate set $\mathcal{S}^t$ of $k = 1e2$ generated subsets from the original set. This is done to calculate the estimated z-score of the point, which is finally used as the target value for our model. Consequentially, if after training the model has converged to a mean low error for the estimated scores, the network parameters are deemed optimized to substitute the Monte Carlo process in estimating daily SS scores for the currently listed option contracts [Kim et al., 2021].

```

1 SS_Score(maturities, candidate_set, p, k, sigma)
2   optimal_parameter = OptimizeForParameter(candidate_set)
3
4   for candidate_point in candidate_set
5     // Generate k subsets of size Length(candidate_set) * p
6     // that contain the candidate_point
7     subset <- GenerateSubset(candidate_set, candidate_point, p, k)
8     suboptimal_parameter <- OptimizeForParameter(subset)
9
10    // For each maturity, calculate the parameter's value using
11    // one of the equations at Eq. 3, 4, 5.
12    error <- Mean(AbsoluteDifference(
13      f(maturities, suboptimal_parameter) - candidate_set
14    ))
15
16    ss_score <- Sqrt(e^(-Error / Sigma))
17    Append ss_score to subset_scores
18
19  z_score <- (ss_score - Mean(ss_score)) / StdDev(ss_score)
20
21  return z_score

```

2.3 Details on the Chosen Encoder Architecture

Unlike recurrent neural networks (RNNs) and convolutional neural networks (CNNs), transformers do not rely on a sequential structure to process data. Instead, they use an attention mechanism, which allows each input element (or token) to consider the relationship with all other elements in the same sequence simultaneously. This is achieved through a series of layers called encoders (for encoding the input data) and decoders (for generating the output).

The self-attention mechanism, which accepts continuous data and calculates attention for different head is computed through value (value), key (key), and query (query) matrices, resulting in an output weighted by the importance of each part of the sequence [Wen et al., 2023, Vaswani et al., 2023]

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (6)$$

Transformers are classified as machine learning algorithms due to their ability to learn patterns and relationships from data. They are trained using supervised techniques. In the context of the specific architecture we implemented, we use only the encoder block, with 4 input dimensions, 8 attention heads per block, and 8 blocks in total. The hidden values are passed to 8 hidden fully-connected (FC) layers, and a final implicit fixed-point iteration layer is used to obtain the scores as SS z-normalized values. The model training stage is carried out on historical option data listed for the **S&P 500** from January to December of 2021. The training epochs adjust the layer weights with an error function that minimizes the difference between its predictions and the Monte Carlo expected z-score values.

2.3.1 Details on the Implicit Layer

3 Methodology and Results

3.1 Data Collection and Pre-Processing

The data used in this work were obtained from historical options data on the *SPY* asset from the US stock exchange, from 2021 to 2022, available on Kaggle. This dataset contains various columns about traded options, indexed by date, including strike prices, future prices, maturities, implied volatilities, among others []. The file can be and was downloaded and stored locally in a specific directory for processing.

3.2 Dataset Description

The dataset contains the following columns relevant to our study, in addition to several others:

- **underlying**: Price of the underlying asset.
- **strike**: Option's strike price.
- **daysToExpiration**: Days remaining until the option's expiration.
- **iv**: Implied volatility.
- **dt**: Record date.

Additionally, we added two calculated columns:

- **maturity**: Calculated by dividing *daysToExpiration* by 360 (working days in a year) to normalize the maturity time.
- **r**: Risk-free rate obtained from the Federal Reserve Economic Data (FRED) for the corresponding period.
- **d**: Dividend yield paid quarterly by the SPY.

3.3 Data Preprocessing

The data preprocessing involved several essential steps, predominantly performed to allow the use of the data directly in the transformer architecture and subsequent use to generate the surface for historical data. The transformer takes four dimensions as input, which are:

1. Time to maturity;
2. The value of the parameter (α , ρ , or ν);
3. **1** if the parameter was calculated for the current day, **0** if it was calculated on the previous day because it does not exist on the current day;
4. Parameters calculated using the previous day's scores.

The target variable is the scores generated using all candidate points for the current day, and the model's error function uses a percentage **p** of points and is trained to minimize the quadratic error function of the scores.

3.4 Results and Surface Comparison

We used a linear interpolation of the daily observed points as the baseline and compared it with the model's performance concerning internal points to a grid generated for points within the daily range, both for maturity values and strike.

Our transformer model was trained on the dataset, with approximately 70 observations, each with a sequence of an average of 700 points. The models were trained for approximately 700 epochs.




Figure 1: Historical loss function from a Transformer training session




Figure 2: Surface generated using a manually optimized SABR model. Note that the non-manual selection of candidate contracts generated a peak in the surface




Figure 3: Comparison of the surface generated with the proposed architecture and the baseline surface. On the left, the SST model surface, on the right, the linearly interpolated surface

4 Conclusion

4.1 Summary

This work presented an approach to adjusting the parameters of the SABR model using a neural network based on the transformer architecture (specifically the *encoder* element). The combination of the SABR model with the transformer architecture was used to replicate the patterns of the daily implied volatility surface, generating a smooth surface without arbitrage opportunities for the existing contracts observed on the day. The fact that the selection and scoring of candidates can be performed automatically, instead of manually, is an extremely efficient solution, especially considering the possibility of using these models on the GPU, which allows for surface reconstruction in an instant, enabling investors to identify trading opportunities in just a few moments, rather than on a daily scale.

4.2 Implications and Future Work

The main focus of this work is indeed the approach using the transformer to replace the manual selection of scores for candidates, paving the way for eventual new research and High-Frequency Trading (HFT) strategies. Future research can explore the integration of additional market factors and improve the chosen architecture, possibly increasing the number of blocks or attention heads. Furthermore, it is important to consider the choice of optimizer for the parameters $\mathbf{P}$, $\mathbf{Q}$, $\mathbf{R}$, as the time required for the preprocessing step was a major limitation during the work. Extending the model to other financial instruments is also a possibility, as it allows for the study of the model's adaptation to different market regimes and economic conditions (both micro- and macro-economic), which may reveal new opportunities for application and development of the proposed method.

References

- M. Ludkovski. Statistical machine learning for quantitative finance. *ANNUAL REVIEW OF STATISTICS AND ITS APPLICATION*, 10:271–295, 2023. ISSN 2326-8298. doi:10.1146/annurev-statistics-032921-042409.
- Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3): 637–654, 1973.
- Bo-Hyun Kim, Daewon Lee, and Jaewook Lee. Local volatility function approximation using reconstructed radial basis function networks. In J Wang, Z Yi, JM Zurada, BL Lu, and H Yin, editors, *ADVANCES IN NEURAL NETWORKS - ISSN 2006, PT 3, PROCEEDINGS*, volume 3973 of *Lecture Notes in Computer Science*, pages 524–530, HEIDELBERGER PLATZ 3, D-14197 BERLIN, GERMANY, 2006. Univ Elect Sci & Technol China; Chinese Univ Hong Kong; Asia Pacific Neural Network Assembly; European Neural Network Soc; IEEE Circuits & Syst Soc; IEEE Computat Intelligence Soc; Int Neural Network Soc; Natl Nat Sci Fdn China; KC Wong Educ Fdn, Hong Kong, SPRINGER-VERLAG BERLIN. ISBN 3-540-34482-9. 3rd International Symposium on Neural Networks (ISSN 2006), Chengdu, PEOPLES R CHINA, MAY 28-31, 2006.
- Cristian Homescu. Implied volatility surface: Construction methodologies and characteristics, 2011.
- Richard Jordan and Charles Tier. Asymptotic approximations to cev and sabr models. *SSRN Electronic Journal*, pages 1–38, May 2011. Available at SSRN: <https://ssrn.com/abstract=1850709> or <http://dx.doi.org/10.2139/ssrn.1850709>.
- Patrick S. Hagan, Deep Kumar, Andrew Lesniewski, and Diana Woodward. Arbitrage-free sabr. *Wilmott*, 2014(69): 60–75, 2014. doi:<https://doi.org/10.1002/wilm.10290>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/wilm.10290>.
- Charles G Broyden. The convergence of a class of double-rank minimization algorithms 1. general considerations. *IMA Journal of Applied Mathematics*, 6(1):76–90, 1970.
- Roger Fletcher. A new approach to variable metric algorithms. *Computer Journal*, 13(3):317–322, 1970.
- Donald Goldfarb. A family of variable-metric methods derived by variational means. *Mathematics of Computation*, 24 (109):23–26, 1970.
- David F Shanno. Conditioning of quasi-newton methods for function minimization. *Mathematics of Computation*, 24 (111):647–656, 1970.
- Hyeonuk Kim, Kyunghyun Park, Junkee Jeon, Changhoon Song, Jungwoo Bae, Yongsik Kim, and Myungjoo Kang. Candidate point selection using a self-attention mechanism for generating a smooth volatility surface under the sabr model. *EXPERT SYSTEMS WITH APPLICATIONS*, 173, JUL 1 2021. ISSN 0957-4174. doi:10.1016/j.eswa.2021.114640.
- Patrick Hagan, Deep Kumar, Andrew Lesniewski, and Diana Woodward. Managing smile risk. *Wilmott Magazine*, 1: 84–108, 01 2002.
- Zico Kolter, David Duvenaud, and Matt Johnson. Deep implicit layers - neural odes, deep equilibrium models, and beyond, 2020. NeurIPS 2020 tutorial. Retrieved from <https://implicit-layers-tutorial.org/>.
- Kenji Kawaguchi. On the theory of implicit deep learning: Global convergence with implicit layers. 2021. URL <https://arxiv.org/abs/2102.07346>.
- Soohan Kim, Seok-Bae Yun, Hyeong-Ohk Bae, Muhyun Lee, and Youngjoon Hong. Physics-informed convolutional transformer for predicting volatility surface. *QUANTITATIVE FINANCE*, 24(2):203–220, JAN 9 2024. ISSN 1469-7688. doi:10.1080/14697688.2023.2294799.
- Patrick Hagan, Andrew Lesniewski, and Diana Woodward. *Probability Distribution in the SABR Model of Stochastic Volatility*, volume 110, pages 1–35. 06 2015. ISBN 978-3-319-11604-4. doi:10.1007/978-3-319-11605-1_1.
- Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David Duvenaud. Neural ordinary differential equations. 2019. URL <https://arxiv.org/abs/1806.07366>.
- Stefano Massaroli, Michael Poli, Jinkyoo Park, Atsushi Yamashita, and Hajime Asama. Dissecting neural odes. 2021. URL <https://arxiv.org/abs/2002.08071>.
- Xuechen Li, Ting-Kam Leonard Wong, Ricky T. Q. Chen, and David Duvenaud. Scalable gradients for stochastic differential equations. 2020. URL <https://arxiv.org/abs/2001.01328>.
- Qingsong Wen, Tian Zhou, Chaoli Zhang, Weiqi Chen, Ziqing Ma, Junchi Yan, and Liang Sun. Transformers in time series: A survey. 2023. URL <https://arxiv.org/abs/2202.07125>.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. 2023. URL <https://arxiv.org/abs/1706.03762>.